



Deep Learning

1

2025-2026 – Pascal Mettes

Lecture 7

Graph Deep Learning

Previous lecture

Lecture	Title	Lecture	Title
1	Intro and history of deep learning	2	AutoDiff
3	Deep learning optimization I	4	Deep learning optimization II
5	Convolutional deep learning	6	Attention-based deep learning
7	Graph deep learning	8	From supervised to unsupervised deep learning
9	Multi-modal deep learning	10	Generative deep learning
11	What doesn't work in deep learning	12	Non-Euclidean deep learning
13	Q&A	14	Deep learning for videos

I have to leave at 10:10 on the dot to be able to make it to a PhD committee at 11:00 in the city center.

This lecture






Graphs

Graph convolutions


Graph attention

Graph applications






Graphs, more common than you think



Many structures are special cases
of graphs




Graphs as geometry



What are graph networks?

The bigger picture:




Notation: $\mathcal{G} = (\mathbf{A}, \mathbf{X})$

- Adjacency matrix $\mathbf{A} \in \mathbb{R}^{N \times N}$
- Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times F}$

Main idea: Pass messages between pairs of nodes & agglomerate

Graph networks

Input: Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times E}$, preprocessed adjacency matrix $\hat{\mathbf{A}}$



$$\mathbf{H}^{(l+1)} = \sigma(\hat{\mathbf{A}}\mathbf{H}^{(l)}\mathbf{W}^{(l)})$$

Node classification:

$$\text{softmax}(\mathbf{z}_n)$$

e.g. Kipf & Welling (ICLR 2017)

Graph classification:

$$\text{softmax}(\sum_n \mathbf{z}_n)$$

e.g. Duvenaud et al. (NIPS 2015)

Link prediction:

$$p(A_{ij}) = \sigma(\mathbf{z}_i^T \mathbf{z}_j)$$


Kipf & Welling (NIPS BDL 2016)

“Graph Auto-Encoders”

1) Graph classification

Make a prediction over the entire graph.

Akin to assigning a label to an entire image.



2) Node classification


Make a prediction for each individual node.

Akin to segmentation for images.




3) Link prediction

Make a prediction for each edge between two nodes.




B




4) Graph generation


Similar in spirit to image/text generation, topic of next week.

Example molecule




Generated molecule






Graphs can be dynamic

Graphs have fixed structures.



But many are subject to




In practice, this change can even be gradual continuous.

Regular structures and graphs


Regular structures are a subset of graphs. I.e., images are grid graphs.



- Convolution + pooling
- Local neighborhood: fixed window
- Constant number of neighbors
- With fixed ordering
- Translation equivariance



- Message passing + coarsening
- Local neighborhood: 1-hop
- Different number of neighbors
- No ordering of neighbors
- Local permutation equivariance

Labelled graph	Degree matrix	Adjacency matrix	Laplacian matrix
	$\begin{pmatrix} 2 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$	$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$	$\begin{pmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{pmatrix}$


Definition of a graph

(in deep learning)

Directed graphs

Vertices $\mathcal{V} = \{1, \dots, n\}$, also called “nodes”

Edges $\mathcal{E} = \{(i, j) : i, j \in \mathcal{V}\} \subseteq \mathcal{V} \times \mathcal{V}$ (directed)




Undirected graphs

Vertices $\mathcal{V} = \{1, \dots, n\}$

Edges $\mathcal{E} = \{(i, j) : i, j \in \mathcal{V}\} \subseteq \mathcal{V} \times \mathcal{V}$ (directed)

Edges $\mathcal{E} = \{\{i, j\} : i, j \in \mathcal{V}\} \subseteq \mathcal{V} \times \mathcal{V}$ (undirected)




Graph neighborhood

The neighborhood of a node consists of all nodes directly connected to it

$$\mathcal{N}(i) = \{j : (i, j) \in \mathcal{E}\}$$

The **degree** of a node is the number of neighbors: $d_i = |\mathcal{N}(i)|$


The diagonal matrix D contains all degrees per node



Attributes

Node features $\mathbf{x}: \mathcal{V} \rightarrow \mathbb{R}^d$, $X = (\mathbf{x}_1, \dots, \mathbf{x}_n)$

Edge features $\mathbf{e}_{ij}: \mathcal{E} \rightarrow \mathbb{R}^{d'}$




Adjacency matrix

An $n \times n$ matrix A , for n nodes

$$A_{ij} = \begin{cases} 1 & \text{if } (i,j) \in \varepsilon \\ 0 & \text{if } (i,j) \notin \varepsilon \end{cases}$$

$(A^z)_{ij}$: number of paths that go from i to j in z steps


	i	j	k	l	m
i		1	1	0	
j					
k	1				
l	0		1		
m					



Adjacency matrix for undirected graphs

The adjacency matrix is symmetric for undirected graphs.


	i	j	k	l	m
i		1	1	0	1
j	1				
k	1				
l	0				
m	1				




Weighted adjacency matrix

When the edges have weights, so does the adjacency matrix.

	i	j	k	l	m
i			w_{ij}	w_{ik}	0
j					
k			w_{ki}		
l			0		
m				w_{mi}	




Final graph input representation



b)

Adjacency matrix, \mathbf{A}
 $N \times N$

	1	2	3	4	5	6
1	■	□	□	■	□	□
2	□	■	□	□	□	□
3	■	■	■	□	□	■
4	■	■	■	■	■	■
5	■	■	■	■	■	■
6	■	■	■	■	■	■




Back to graph networks

We can do a lot of processing on this data structure.

But the pre-defined features are raw inputs.

Graph networks do 1 thing: transform the feature vector per node over layers.


The bigger picture:



Notation: $\mathcal{G} = (\mathbf{A}, \mathbf{X})$

- Adjacency matrix $\mathbf{A} \in \mathbb{R}^{N \times N}$
- Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times F}$



Three perspectives to graph networks



Three “flavours” of GNNs, left-to-right: convolutional, attentional, and general nonlinear message passing flavours. All are forms of message passing. Figure adapted from P. Veličković.

Graph layer as a convolution layer

**Single CNN layer
with 3x3 filter:**



Update for a single pixel:

- Transform messages individually $\mathbf{W}_i \mathbf{h}_i$
- Add everything up $\sum_i \mathbf{W}_i \mathbf{h}_i$

$\mathbf{h}_i \in \mathbb{R}^F$ are (hidden layer) activations of a pixel/node

Full update:



$$\mathbf{h}_4^{(l+1)} = \sigma \left(\mathbf{W}_0^{(l)} \mathbf{h}_0^{(l)} + \mathbf{W}_1^{(l)} \mathbf{h}_1^{(l)} + \dots + \mathbf{W}_8^{(l)} \mathbf{h}_8^{(l)} \right)$$

Which assumptions from images are no longer valid?

Number of neighbors per node no longer fixed.



No more ordering between neighbours.

Extending convolutions to graphs



Convolution in CNNs

CNNs perform localized convolutions. Neighbours participating in the convolution at the center pixel are highlighted in gray.




Localized Convolution in GNNs


GNNs can perform localized convolutions mimicking CNNs. Hover over a node to see its immediate neighbourhood highlighted on the left. The structure of this neighbourhood changes from node to node.

Graph convolution layer

Consider this
undirected graph:



Calculate update
for node in red:




Stacking graph convolution layers

Each layer aggregates information from their direct neighbors.


At the end of each layer, we add a non-linearity such as a ReLU.

We can increase complexity and receptive field simply by stacking multiple layers.

Consider this undirected graph:



Calculate update for node in red:



Update rule:
$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

Graph convolution layer in matrix form

$$f(X, A) := \sigma \left(\mathbf{D}^{-1/2} (\mathbf{A} + \mathbf{I}) \mathbf{D}^{-1/2} \mathbf{X} \mathbf{W} \right)$$

$A \in \mathbb{R}^{n \times n}$:= The adjacency matrix

$I \in \mathbb{R}^{n \times n}$:= The identity matrix

$D \in \mathbb{R}^{n \times n}$:= The degree matrix of $A + I$

$X \in \mathbb{R}^{n \times d}$:= The input data (i.e., the per-node feature vectors)

$W \in \mathbb{R}^{d \times w}$:= The layer's weights

$\sigma(\cdot)$:= The activation function (e.g., ReLU)

Let's break it down

$$f(\mathbf{X}, \mathbf{A}) := \sigma\left(\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW}\right)$$

The diagram illustrates the components of the equation $f(\mathbf{X}, \mathbf{A})$. It shows a large bracket under the term $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}$ labeled "Normalize adjacency matrix". Below this, a bracket under the term \mathbf{XW} is labeled "Aggregate". The entire expression $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW}$ is enclosed in a large bracket labeled "Update". Within the "Update" bracket, a bracket under the term $\mathbf{A} + \mathbf{I}$ is labeled "Add self-loops".

Let's break it down

Add ones to diagonal, needed because each node should pass its own vector through.

$$f(\mathbf{X}, \mathbf{A}) := \sigma\left(\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW}\right)$$

The diagram illustrates the components of the formula:

- Add self-loops:** A red box highlights the term $\mathbf{A} + \mathbf{I}$, indicating the addition of identity matrices to the adjacency matrix.
- Normalize adjacency matrix:** Brackets group the first two terms, $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})$, representing the normalization of the adjacency matrix by the degree matrix.
- Aggregate:** Brackets group the third term, $\mathbf{D}^{-1/2}\mathbf{XW}$, representing the aggregation of node features.
- Update:** Brackets group the entire right-hand side of the equation, representing the final updated node representation.

Let's break it down

This step essentially normalizes the adjacency matrix. I will show how in a few slides.

$$f(\mathbf{X}, \mathbf{A}) := \sigma\left(\boxed{\mathbf{D}^{-1/2}} (\mathbf{A} + \mathbf{I}) \boxed{\mathbf{D}^{-1/2}} \mathbf{XW}\right)$$

The diagram illustrates the components of the function $f(\mathbf{X}, \mathbf{A})$ through a series of nested brackets:

- The innermost bracket, enclosed in red boxes, contains $\mathbf{D}^{-1/2}$, $(\mathbf{A} + \mathbf{I})$, and $\mathbf{D}^{-1/2}$.
- A bracket below this labeled "Add self-loops" covers the term $(\mathbf{A} + \mathbf{I})$.
- A bracket below that labeled "Normalize adjacency matrix" covers the entire product $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}$.
- A bracket below that labeled "Aggregate" covers the product $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW}$.
- The outermost bracket, enclosed in red boxes, covers the entire expression $\sigma(\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW})$.
- A bracket at the bottom labeled "Update" covers the entire function $f(\mathbf{X}, \mathbf{A})$.

Let's break it down

Just a standard linear layer and a non-linearity.

$$f(\mathbf{X}, \mathbf{A}) := \sigma(\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW})$$


The diagram illustrates the components of the function $f(\mathbf{X}, \mathbf{A})$ through a series of nested brackets:

- A bracket labeled "Normalize adjacency matrix" encloses the term $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}$.
- A bracket labeled "Add self-loops" is positioned above the "Normalize adjacency matrix" bracket, covering the term $\mathbf{A} + \mathbf{I}$.
- A bracket labeled "Aggregate" is positioned below the "Normalize adjacency matrix" bracket, covering the term \mathbf{XW} .
- A bracket labeled "Update" is positioned at the bottom, enclosing the entire expression $\sigma(\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW})$.

Two specific terms, σ and \mathbf{XW} , are highlighted with red boxes.

Let's break it down

Just a standard linear layer and a non-linearity.

$$f(\mathbf{X}, \mathbf{A}) := \sigma(\boxed{\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW}})$$


The diagram illustrates the components of the function $f(\mathbf{X}, \mathbf{A})$:

- A red box highlights the term $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}$.
- A dashed box encloses the term $\mathbf{A} + \mathbf{I}$, with a bracket labeled "Add self-loops".
- To the right of the dashed box, another dashed box encloses the term \mathbf{XW} , with a bracket labeled "Normalize adjacency matrix".
- A bracket labeled "Aggregate" spans the entire term $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW}$.
- A bracket labeled "Update" spans the entire expression $\sigma(\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW})$.

Rewriting into 2 steps

$$\tilde{\mathbf{A}} := \mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}$$

$$\tilde{\mathbf{A}}_{i,j} := \begin{cases} \frac{1}{\sqrt{d_{i,i}d_{j,j}}}, & \text{if there is an edge between node } i \text{ and } j \\ 0, & \text{otherwise} \end{cases}$$


$$f(\mathbf{X}, \mathbf{A}) := \sigma(\tilde{\mathbf{A}}\mathbf{X}\mathbf{W})$$

$$\mathbf{D} := \begin{bmatrix} d_{1,1} & 0 & 0 & \dots & 0 \\ 0 & d_{2,2} & 0 & \dots & 0 \\ 0 & 0 & d_{3,3} & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & d_{n,n} \end{bmatrix}$$

$$\mathbf{D}^{-1/2} := \begin{bmatrix} \frac{1}{\sqrt{d_{1,1}}} & 0 & 0 & \dots & 0 \\ 0 & \frac{1}{\sqrt{d_{2,2}}} & 0 & \dots & 0 \\ 0 & 0 & \frac{1}{\sqrt{d_{3,3}}} & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & \frac{1}{\sqrt{d_{n,n}}} \end{bmatrix}$$

$$f(X, \mathbf{A}) := \sigma(\tilde{\mathbf{A}}\mathbf{X}\mathbf{W})$$


Left side of the equation



$$\begin{aligned} \bar{\mathbf{x}}_i &= \sum_{j=1}^n \tilde{a}_{ij} \mathbf{x}_j \\ &= \sum_{j \in \text{Neigh}(i)} \tilde{a}_{ij} \mathbf{x}_j \\ &= \sum_{j \in \text{Neigh}(i)} \frac{1}{\sqrt{d_{i,i} d_{j,j}}} \mathbf{x}_j \end{aligned}$$

$$f(X, A) := \sigma(\tilde{A}XW)$$

Right side of the equation



Why add a normalization step?

Do we even need it? Let's see what happens without it:

$$\hat{\mathbf{A}} := \mathbf{A} + \mathbf{I}$$

Normalization dropped,
only self-loop
retrained

$$f_{\text{unnormalized}}(\mathbf{X}, \mathbf{A}) := \sigma(\hat{\mathbf{A}}\mathbf{X}\mathbf{W})$$

Layer update remains the
same

$$\bar{x}_i = \sum_{j=1}^n \hat{a}_{i,j} x_j$$


$$= \sum_{j=1}^n \mathbb{I}(j \in \text{Neigh}(i)) x_j$$

$$= \sum_{j \in \text{Neigh}(i)} x_j$$



Problem! More neighbors =
bigger sum.

Huge bias when training graph
networks

Why not simply divide by the node degree?



Visualizing node representations



Alternative: graph layer as attention

Similar but including attention as aggregation: $y_i = h\left(\sum_{j \in \mathcal{N}(i)} a_{ij} \mathbf{z}_j\right)$

Using self-attention:


$$a_{ij} = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}(i)} \exp(e_{kj})},$$

where e_{ij} are the self-attention weights (like query == key)

$$e_{ij} = \text{LeakyReLU}([\mathbf{x}_i \mathbf{W}, \mathbf{x}_j \mathbf{W}] \cdot \mathbf{u})$$

\mathbf{u} is a weight vector.

The four steps of a graph attention layer



$$z_i^{(l)} = W^{(l)} h_i^{(l)}, \quad (1)$$

$$e_{ij}^{(l)} = \text{LeakyReLU}(\vec{a}^{(l)^T} (z_i^{(l)} || z_j^{(l)})), \quad (2)$$

$$\alpha_{ij}^{(l)} = \frac{\exp(e_{ij}^{(l)})}{\sum_{k \in \mathcal{N}(i)} \exp(e_{ik}^{(l)})}, \quad (3)$$

$$h_i^{(l+1)} = \sigma \left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij}^{(l)} z_j^{(l)} \right), \quad (4)$$

Connecting graphs, convolutions, and transformers


Transformers operate on a complete graph (adjacency matrix with all 1's).

With attention-based GCN, we recover the Transformer.

Architecture	Domain Ω	Symmetry group \mathfrak{G}
CNN	Grid	Translation
<i>Spherical CNN</i>	Sphere / SO(3)	Rotation SO(3)
<i>Intrinsic / Mesh CNN</i>	Manifold	Isometry $\text{Iso}(\Omega)$ / Gauge symmetry SO(2)
GNN	Graph	Permutation Σ_n
<i>Deep Sets</i>	Set	Permutation Σ_n
<i>Transformer</i>	Complete Graph	Permutation Σ_n
LSTM	1D Grid	Time warping

Optimizing graph networks

Input: Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times E}$, preprocessed adjacency matrix $\hat{\mathbf{A}}$



$$\mathbf{H}^{(l+1)} = \sigma(\hat{\mathbf{A}}\mathbf{H}^{(l)}\mathbf{W}^{(l)})$$

Node classification:

$$\text{softmax}(\mathbf{z}_n)$$

e.g. Kipf & Welling (ICLR 2017)

Graph classification:

$$\text{softmax}(\sum_n \mathbf{z}_n)$$

e.g. Duvenaud et al. (NIPS 2015)

Link prediction:

$$p(A_{ij}) = \sigma(\mathbf{z}_i^T \mathbf{z}_j)$$


Kipf & Welling (NIPS BDL 2016)

“Graph Auto-Encoders”

Pooling in graph networks



Specifically for graph classification, pooling is an optional operators.

Pool nodes together to save compute, requires updating the adjacency matrix.





Applications of graph networks



Drug discovery




Modeling the spread of deceases





Social networks



Recommendation





Traffic forecasting





The model architecture for determining optimal routes and their travel time.


Scene graph generation of visual data



Point cloud classification




Object interactions



(i) Image

(ii) HOI candidates

(iii) HOI result



(iv) Initial
HOI graph


(v) Parse
graph learning

(vi) Message
passing


(vii) Final
parse graph

Joint Inference

Text classification



Particle physics



Next lecture

Lecture	Title	Lecture	Title
1	Intro and history of deep learning	2	AutoDiff
3	Deep learning optimization I	4	Deep learning optimization II
5	Convolutional deep learning	6	Attention-based deep learning
7	Graph deep learning	8	From supervised to unsupervised deep learning
9	Multi-modal deep learning	10	Generative deep learning
11	What doesn't work in deep learning	12	Non-Euclidean deep learning
13	Q&A	14	Deep learning for videos

Learning and reflection

Understanding Deep Learning: Chapter 13

Thank you !