


Exercises

1	2	3	4	5	6	7	8	9	10
11	12	13	14						

Surname, First name

Machine Learning 1 (52041MAL6Y)

Final Exam

*Final exam for Machine Learning 1 - 26 October 2021*

Pay attention to the following

- **Write your name and student number on the front page** (don't forget to mark the digits)
- Write all your answers on this exam booklet; it will be scanned and digitally graded.
- Write your **answers inside the boxes** (it is ok, to slightly go over the margins though).
- The answer boxes should be large enough for your answer
- **If you need to empty the answer box in order to start over, ask the invigilator for a blank sticker**

During the exam you are allowed to use:

- One double-sided handwritten cheat sheet.
- A calculator (though not strictly necessary)

About the multiple choice questions:

- You should fill the boxes and not just check them. E.g.

Correct marking:

Incorrect marking:

- In case you want to correct your answer, clearly indicate this (e.g. by filling all boxes and use another indicator such as an arrow to indicate your choice). We can then resolve this during grading.
- **The grading of the multiple choice questions is based on all or nothing:** you get full points for ticking off the correct boxes and leaving the others blank, but get zero points if you didn't check all correct answers, or ticked off an incorrect answer.

About the open questions:

- We work with a partial grading: you can get points even if you don't manage to solve the full question.

Please scan over the questions before you start to get an impression of the content

- In total **50 points** can be earned (+ some bonus) divided over the following 5 categories.
- You have **3 hours** for the exam (except for those with pre-approved extensions)
- Exercises 1-10: Multiple choice (10.0 pts)
- Exercise 11: Probability theory (5.5 pts + 1.0 bonus)
- Exercise 12 Discrete latent variable modeling (14.0 pts + 1.5 bonus)
- Exercise 13 Continuous latent variable modeling (5.5 pts + 1.0 bonus)
- Exercise 14 Support Vector Machines (15.0 pts)



ML vs MAP

- 1p 1 Given a dataset $\mathcal{D} = \{x_n\}_{n=1}^N$. The data is normally distributed $\mathcal{N}(x_n|\mu, \sigma^2)$ and we assume a Gaussian prior over $\mu : \mathcal{N}(\mu|0, \sigma_0^2)$. Furthermore, the variance σ^2 is assumed to be known. Let μ_{ML} and μ_{MAP} respectively be the ML and MAP estimates for μ . How does $|\mu_{ML} - \mu_{MAP}|$ change as (i) $\sigma_0 \rightarrow 0$, (ii) $\sigma_0 \rightarrow \infty$, (iii) $N \rightarrow \infty$.
- (a) (i) decrease (ii) increase (iii) decrease.
 - (b) (i) decrease (ii) increase (iii) increase.
 - (c) (i) increase (ii) decrease (iii) decrease.
 - (d) (i) increase (ii) decrease (iii) increase.

Changing decision boundaries

- 1p 2 Consider a binary classification problem. Suppose I have trained a model on a linearly separable training set, and now I get a new labeled data point which is correctly classified by the model, and far away from the decision boundary. If I now add this new point to my earlier training set and re-train, in which cases is the learned decision boundary likely to change?
- When my model is a perceptron.
 - When my model is logistic regression.
 - When my model is an SVM.
 - When my model is Gaussian discriminant analysis.

Misclassification error

- 1p 3 After fitting a Logistic Regression model with two classes to the training data we calculate the confusion matrix of the model on the test data with N observations: $\begin{pmatrix} A & B \\ C & D \end{pmatrix}$. Which of the following formulas could you use to estimate the misclassification error:
- $\frac{1}{N}(A + D)$.
 - $1 - \frac{1}{N}(A + D)$.
 - $\frac{1}{N}(B + C)$.
 - $1 - \frac{1}{N}(B + C)$.



Probabilistic models

1p **4** In classification, there are three approaches, discriminant functions, probabilistic generative models and probabilistic discriminative models. The following are statements about probabilistic generative models and probabilistic discriminative models. Which of the following statements is true?

- Logistic regression is a probabilistic discriminative model.
- In probabilistic discriminative models, the posterior class probabilities $p(C|x)$ are modeled directly.
- In probabilistic discriminative models, the prior probability of class $p(C)$ is modeled.
- In generative models, the class conditional probability $p(x|C)$ are modeled directly.

Neural networks: representation power

1p **5** Consider a neural network with L layers, and H hidden units per hidden layer, and O output units.

Given a fixed number of network parameters (say fix the number L, H) and activation functions

- $h(x) = \max(0, x)$ for the hidden units, one most effectively increases the network complexity by increasing L rather than H .
- A 5 layer neural network with activation functions $h(x) = \frac{1}{1+e^{-x}}$ for the hidden layers is able to approximate any continuous function with compact support to arbitrary precision via suitable choice for H .
- A 5 layer neural network with activation functions $h(x) = x^3$ for the hidden layers is able to approximate any continuous function with compact support to arbitrary precision via suitable choice for H .
- With the same amount of output units O , a linear regression model is at least as expressive as a deep neural network with $L = 100$ layers, any choice of H , and activation functions $h(x) = x$.

The kernel trick

1p **6** What is the advantage of using kernels in methods such as support vector machines?


- They can simulate an infinite dimensional feature space.
- They will always reduce the number of support vectors.
- They reduce the risk of getting stuck in local minima.
- They make it possible to model non-linear decision boundaries.

SVM and underfitting

- 1p 7 Suppose you are given the following binary dataset and trained a SVM that solves

$$\underset{\mathbf{w}, \mathbf{b}, \{\xi_n\}}{\text{minimize}} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n \quad \text{subject to} \quad \begin{array}{l} \forall_{n=1,\dots,N}: t_n y_n \geq 1 - \xi_n \\ \forall_{n=1,\dots,N}: \xi_n \geq 0 \end{array}$$

using a Gaussian kernel $k(\mathbf{x}, \mathbf{x}') = \exp(-\frac{1}{2\sigma^2} \|\mathbf{x} - \mathbf{x}'\|^2)$ that gave the following decision boundary:



You suspect that the SVM is underfitting your dataset. Should you try to increase or decrease the C parameter? Increase or decrease σ^2 ?

- a Increase C, decrease σ^2
- b Increase C, increase σ^2
- c Decrease C, increase σ^2
- d Decrease C, decrease σ^2

Local vs global optimality

- 1p 8 In which of the following can we expect to find the globally optimal solution for the model parameters? In all cases assume i.i.d. assumption on the dataset $D = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\}$ and assume all (undefined) parameters/functions to be known.

• minimize $\frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$ subject to $\begin{array}{l} \forall_{n=1,\dots,N}: t_n y_n \geq 1 - \xi_n \\ \forall_{n=1,\dots,N}: \xi_n \geq 0 \end{array}$

with $y_n = \mathbf{w}^T \phi(\mathbf{x}_n) + \mathbf{b}$ and with $\phi(\mathbf{x})$ some non-linear feature transformation of input \mathbf{x} .

• maximize $\underset{\mathbf{w}, \mathbf{b}, \beta}{p(\mathbf{w}|D)} = \frac{p(D|\mathbf{w})p(\mathbf{w})}{p(D)}$

using Gaussian predictive distribution $p(t|\mathbf{x}, \mathbf{w}) = \mathcal{N}(t|\mathbf{w}^T \mathbf{x} + \mathbf{b}, \beta^{-1})$ and prior $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\boldsymbol{\mu}, \mathbf{S})$.

• maximize $\underset{\{\pi_k\}, \{\boldsymbol{\mu}_k\}, \sigma}{p(D|\mathbf{w}, \{\pi_k\}, \{\boldsymbol{\mu}_k\}, \sigma)}$

with $p(\mathbf{x}) = \sum_{k=1}^5 \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \sigma^2 \mathbf{I}) \pi_k$, with $\sum_{l=1}^5 \pi_l = 1$.

• minimize $\underset{\mathbf{w}}{\sum_{n=1}^N (NN_{\mathbf{w}}(\mathbf{x}_n) - t_n)^2}$ using Gradient Descent

with $NN_{\mathbf{w}}$ a deep neural network with ReLU activations and parametrized by weights \mathbf{w} .

• minimize $\underset{\mathbf{w}}{\sum_{n=1}^N (NN_{\mathbf{w}}(\mathbf{x}_n) - t_n)^2}$ using Stochastic Gradient Descent

with $NN_{\mathbf{w}}$ a deep neural network with ReLU activations and parametrized by weights \mathbf{w} .

Gaussian process and its kernel

- 1p 9 The values $f(\mathbf{x}_i)$ of a Gaussian process f evaluated at a fixed set of N points $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, with each $\mathbf{x}_i \in \mathbb{R}^D$ can be described by a multi-variate Gaussian. Which of the following statements are true?

- In order to compute the covariance matrix of the multi-variate Gaussian one needs to evaluate the kernel for all data point pairs.
- The kernel makes its appearance in both the covariance matrix as well as the mean of the multi-variate Gaussian.
- The kernel describes how two function values $f(\mathbf{x})$ and $f(\mathbf{y})$ covariate for any choice of $\mathbf{x}, \mathbf{y} \in \mathbb{R}^D$.
- The covariance matrix of the multi-variate Gaussian is a $D \times D$ matrix.




SVM: decision boundary

- 1p 10 Consider the following two-class SVM optimization problem with data set $\mathcal{D} = \{(\mathbf{x}_n, t_n)\}_{n=1}^N$ with $\mathbf{x}_n \in \mathbb{R}^2$ and $t_n \in \{-1, 1\}$:

$$\underset{\mathbf{w}, \mathbf{b}, \{\xi_n\}}{\text{minimize}} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n \quad \text{subject to} \quad \begin{array}{ll} \forall_{n=1,\dots,N}: & t_n(\mathbf{w}^T \mathbf{x}_n - \mathbf{b}) \geq 1 - \xi_n \\ \forall_{n=1,\dots,N}: & \xi_n \geq 0 \end{array}$$

The data set is depicted in the figure below. The 11 gray dots correspond to data points with labels $t_n = +1$, and the 9 black stars to data points with $t_n = -1$.



Instead of solving the above problem directly you consider solving it using the kernel trick and try out 2 kernels: namely $k_A(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$ and $k_B(\mathbf{x}, \mathbf{x}') = \exp(-\frac{1}{2\sigma^2} \|\mathbf{x} - \mathbf{x}'\|^2)$ with some choice for σ .

Which of the following statements are correct?

- The solid line (decision bounday I) could represent the learned decision boundary using k_A and with the setting $C \rightarrow \infty$.
- The solid line (decision bounday I) could represent the learned decision boundary for kernel k_A and with the setting $C = 1$.
- The dashed line (decision bounday II) could represent the learned decision boundary for kernel k_B and with the setting $C \rightarrow \infty$.
- The dashed line (decision bounday II) could represent the learned decision boundary for kernel k_B and with the setting $C = 1$.



ML1 student on a date (Probability theory)

A friend set you up for date with someone called Alex. You are afraid of getting in a sequence of dates going nowhere, so you decide to adopt a Bayesian approach to dating; you will setup a probabilistic model to help you decide whether or not a second date would be worthwhile.

You heard that successful couples have a shared taste in music, so you focus on this feature. You pick your favorite album and assume that, if Alex were indeed to be your soulmate, then there's an 80% chance of you both liking the album. You also assume that you have a very sophisticated taste and that only 1 in 1,000 people like this album. So you reason, there is still a 1/1,000 chance that Alex likes the album even if you don't turn out to be soulmates. Finally, you guess that the chance for finding a match is 1% in the first place.

You stick to the following conventions:

- Define M to be the binary random variable of finding a match ($M = 1$) with the date.
 - Let A be the random variable for someone liking the album ($A = 1$) or not ($A = 0$).
 - You decide for a second date if the probability of Alex being a match is larger than 25%.
 - Your definition of an *enjoyable evening* is one where you can talk all night with a likeminded person about your favorite album.

1.5p **11a** What is the probability of you having an *enjoyable evening* with Alex?

Give your answer in % in one decimal precision (xx.x%) or provide it as a fraction.

2p **11b** So you go out with Alex and things are going smoothly. You just discovered that Alex really likes your favorite album too! 😍 Wow, did you find your soulmate? Compute the probability for a match! Give your answer in % in one decimal precision (xx.x%) or provide it as a fraction.

- 1p **11c [BONUS]** Alex tells you that the album is in all the "all-time-favorite" album charts and it is actually not that special that you both like it; around 1 in 50 people like it. Your heart sinks... 😱 This greatly reduces the probability of being soulmates. Recompute it by updating your initial assumption for the probability for $A = 1$ for non-soulmates ($M = 0$) with the new fact. Will you go for a second date?

 - (a) Yes, few it's just above 25%!
 - (b) No way, not with less than 25% chance of Alex being your soulmate!

2p **11d** "Wait! Are you taking notes on a date!?!?" Alex asks. Whoops, awkward.... 😅 You explain you were computing the posterior for having just found your soulmate". "You did what!?", Alex says, "You should have told me, I'm quite the Bayesian myself you know! With friends were building such a soulmate prediction model based on our collective experiences!"

Alex explains their model is parameterized by weights w and predicts the probability of a match M given some experience vector x . They decided on a prior for w and update their choice for w given a collected dataset $D_N = \{x_i\}_{i=1}^N$ of N experiences. They recompute the most probable model parameters, given the data, every time a new experience x_{N+1} is added (creating D_{N+1}). "The dataset is getting too large and it is increasingly demanding to compute the most probable w . Super annoying..." Alex says.

OK, you got this! Propose to Alex... a solution to this problem!

Discrete latent variable modeling

We have access to a dataset $D = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ of N observations which we assume to be independently and identically distributed and generated by a process that depends on a discrete latent variable $z \in \{C_1, \dots, C_K\}$. I.e., we assume $\mathbf{x} \sim p(\mathbf{x}|z)$, K discrete latent classes, and that $z \sim p(z)$. We will denote with \mathbf{z} the one-hot encoding of z . I.e., $\mathbf{z} = (z_1, \dots, z_k)^T \in \{0, 1\}^K$ where $z_k = 1$ if $z = C_k$ and 0 otherwise.

We thus believe that each datapoint x_i is associated with a latent z , however, there is no way of observing it. So, instead we try to infer the latent variables using a generative model. We will model the latent conditionals with a normal distribution with parameters σ and μ_k for each class via

$$p(\mathbf{x}|z = C_k, \boldsymbol{\mu}_k, \sigma) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \mathbf{I}\sigma^2).$$

We will model the prior on the latent with a Bernoulli distribution

$$p(\mathbf{z}|\{\pi_k\}) = \prod_{k=1}^K \pi_k^{z_k},$$

with model parameters $\{\pi_k\}_{k=1}^K$ that satisfy $\sum_{k=1}^K \pi_k = 1$. We want to derive the optimal model parameters $\{\pi_k\}$, $\{\mu_k\}$ and σ , and intend to do so via the Expectation Maximization (EM) algorithm.

- 1p **12a** Suppose we have fully optimized the model and found the optimal parameters $\{\pi_k\}$, $\{\mu_k\}$ and σ . How could we generate/sample new data points with this model?

- 2.5p **12b** Give an expression for the marginal $p(\mathbf{x}|\{\pi_k\}, \{\boldsymbol{\mu}_k\}, \sigma)$ and derive the likelihood $p(D|\{\pi_k\}, \{\boldsymbol{\mu}_k\}, \sigma)$ for D being generated by our model. Please mention if you relied on any assumptions.

2p **12c** Write down the expression for the posterior latent class probability $p(z = C_k | \mathbf{x}, \{\pi_k\}, \{\boldsymbol{\mu}_k\}, \sigma)$.

2.5p **12d** Let us refer to the posteriors as responsibilities, denoted with $r_{nk} = p(z = C_k | \mathbf{x}_n, \{\pi_k\}, \{\boldsymbol{\mu}_k\}, \sigma)$. We then maximize the log-likelihood with respect to parameters $\{\pi_k\}$, $\{\boldsymbol{\mu}_k\}$, and σ , while keeping the responsibilities r_{nk} fixed. For example, the update rule for the σ in terms of r_{nk} then becomes

$$\sigma^2 = \frac{1}{N} \frac{1}{d} \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|\mathbf{x} - \boldsymbol{\mu}_k\|^2.$$

Derive the update rule for $\boldsymbol{\mu}_k$ in terms of r_{nk} and explain in words what this update rule does.

Hint: Make use of the fact that $\frac{\partial}{\partial \boldsymbol{\mu}_k} \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \sigma^2) = \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \sigma^2) \frac{1}{\sigma^2} (\mathbf{x} - \boldsymbol{\mu}_k)^T$.

3p **12e** Derive the update rule for π_k in terms of r_{nk} and explain in words what this update rule does.

1p **12f** Describe the EM algorithm and how the update rules are used in it.

1p **12g** Will the EM algorithm converge to a globally optimal solution?

- a Yes b No

- 1p **12h** The K -means clustering algorithm shows strong resemblance to the EM algorithm. In fact, we can turn the above EM algorithm into K -means by directly modeling the posterior probabilities r_{nk} with a function $q(z|x)$, instead of modeling it with $p(z|x, \{\pi_k\}, \{\mu_k\}, \sigma)$ using $p(x|z = C_k, \mu_k, \sigma)$ and $p(z, \{\pi_k\})$. We will refer to this function as the encoder. How should you define $q(z|x)$ to obtain K -means?

Write down the expression for the encoder (i.e. write $q(z|x) = \dots$) or describe in words what this function does. Indicate on which of the parameters $\{\pi_k\}$, $\{\mu_k\}$ and/or σ it depends.

- 1.5p **12i [BONUS]** K -means clustering suffers from the fact that the cluster sizes are equal for each latent class. Can you think of an adaptation based on the above that allows to handle clusters with different sizes? What changes to the encoder and/or variable update rules should be made?

Continuous latent variable modeling

We now consider a dataset $D = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ of observations $\mathbf{x}_n \in \mathbb{R}^d$ which you assume come from a multi-variate normal distribution $p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \mathbf{S})$. You assume that the observations are deterministically associated with a continuous latent variable $\mathbf{z} \in \mathbb{R}^M$ via the model

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}.$$

- 2.5p **13a** Through the above model, the latent \mathbf{z} is also a random variable. What is the expectation and covariance of \mathbf{z} ? Compute $\mathbb{E}[\mathbf{z}]$ and $\text{Cov}[\mathbf{z}]$.

- 1p **13b** Provide an expression for the distribution $p(\mathbf{z})$ in terms \mathbf{W} , \mathbf{b} , μ and \mathbf{S} .

- 1p **13c** The individual elements in the latent vector $\mathbf{z} = (z_1, \dots, z_M)^T \in \mathbb{R}^M$ will generally be correlated. What if we were interested in finding a \mathbf{W} such that \mathbf{z} will consist of uncorrelated elements, what would be a suitable algorithm/method for obtaining such a \mathbf{W} ?

- 1p **13d** Suppose we only have access to a sampling algorithm that can sample from normal distributions $N(\mathbf{0}, \mathbf{I})$ with zero mean and unit diagonal covariance matrix. Which trick could you use to still sample variables that are distributed according to $p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \mathbf{S})$ for arbitrary $\boldsymbol{\mu}$ and \mathbf{S} ? Describe the steps.

- 1p **13e [BONUS]** Proof that with your choice for \mathbf{W} in exercise (c) the elements z_i are indeed uncorrelated.

SVM regression




Figure 1: (a) The tolerance error function E_ϵ . (b) Illustration of SVM regression.

Consider the problem of fitting a function $y(\mathbf{x}; \mathbf{w}, \mathbf{b}) = \mathbf{w}^T \phi(\mathbf{x}) + \mathbf{b}$ to a dataset of point pairs $D = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\}$, with $\mathbf{x}_n \in \mathbb{R}^d$ and $t_n \in \mathbb{R}$. For notational convenience we drop the dependency on parameters \mathbf{w} and \mathbf{b} of the model and simply write $y(\mathbf{x})$. Furthermore, we may simply denote $y(\mathbf{x}_n)$, the prediction for the n^{th} data point, with y_n . We fit by solving the following problem

$$\min_{\mathbf{w}, \mathbf{b}} C \sum_{n=1}^N E(y_n, t_n) + \frac{1}{2} \|\mathbf{w}\|^2, \quad (1)$$

in which $E(y(\mathbf{x}), t)$ penalizes errors made by the model. Instead of using the usual squared error loss (see Figure 1(a)), we allow for a tolerance of ϵ . This is done by giving a 0 penalty if the prediction lies within the tolerance interval, and linearly penalize proportional to the distance towards the tolerance region using the following function (see also Figure 1(a))

$$E_\epsilon(y(\mathbf{x}), t) = \begin{cases} 0 & \text{if } |y(\mathbf{x}) - t| < \epsilon \\ |y(\mathbf{x}) - t| - \epsilon & \text{otherwise} \end{cases}.$$

The penalty-free region is then given by $y(\mathbf{x}) - \epsilon \leq t_n \leq y(\mathbf{x}_n) + \epsilon$ and will be referred to as the ϵ -tube.

Akin to the soft-margin support vector machine case, we allow for some slack and adjust the tolerance interval by possibly adding $\xi_n \geq 0$ at the top of the interval, or subtract $\hat{\xi}_n \geq 0$ at the bottom of the interval. We then obtain an equivalent inequality constraint optimization problem:

$$\min_{\mathbf{w}, \mathbf{b}, \{\xi_n\}, \{\hat{\xi}_n\}} C \sum_{n=1}^N (\xi_n + \hat{\xi}_n) + \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{subject to} \quad \begin{aligned} \xi_n &\geq 0, & (I) \\ \hat{\xi}_n &\geq 0, & (II) \\ t_n &\leq y_n + \epsilon + \xi_n, & (III) \\ t_n &\geq y_n - \epsilon - \hat{\xi}_n. & (IV) \end{aligned}$$

Note that this is equivalent to (1) as the slack variables ξ_n and $\hat{\xi}_n$ capture the error beyond the tolerance ϵ . This problem is visualized in Figure 1b.

- 2p **14a** Write down the primal Lagrangian for this inequality constrained optimization problem. Use the following symbols for the Lagrange multipliers: for constraint (I) use μ_n , for (II) use $\hat{\mu}$, for (III) use a_n , and for (IV) use \hat{a}_n .

- 3p **14b** In the lectures we considered the stationarity conditions separately from the KKT conditions; in this exam we do the same. Write down the KKT conditions (not including the stationarity conditions).

1p **14c** How many KKT conditions do we have?

2p **14d** In order to obtain the dual Lagrangian we need to compute the stationarity conditions. When optimizing the Lagrangian with respect to primal variables ξ_n and $\hat{\xi}_n$ respectively we obtain the following extra constraints on a_n and \hat{a}_n

$$\begin{aligned} \text{for all } n = 1, \dots, N : \quad a_n &= C - \mu_n, \\ \text{for all } n = 1, \dots, N : \quad \hat{a}_n &= C - \hat{\mu}_n. \end{aligned} \tag{2}$$


Complete the set of stationarity conditions by deriving them for the primal variables w and b.

- 1p **14e** The new constraints for the dual variables a_n and \hat{a}_n (see Eq. (2) in exercise (d)), together with the original KKT conditions, define box constraints (min-max intervals) for a_n and \hat{a}_n . Write them down.


- 3p **14f** Without actually solving the dual problem, we can already tell that there will be different type of solutions for the values of the dual variables. In the following three figures, we ask you to encircle the requested instances:

1. In subfigure (i) encircle the points for which both $a_n = 0$ and $\hat{a}_n = 0$.
2. In subfigure (ii) encircle the points for which $0 < a_n < C$.
3. In subfigure (iii) encircle the points for which $\hat{a}_n = C$.


Hint: The KKT and optimality conditions allow you to reason about the solution.



Encircle all points for which $a_n = 0$ and $\hat{a}_n = 0$



Encircle all points for which $0 < a_n < C$



Encircle all points for which $\hat{a}_n = C$

Encircle the points that satisfy the specified criteria.

- 1p **14g** The stationarity conditions that you computed in question (d) should give you an expression for w in terms of the dual variables. Use it to derive a dual formulation for the predictive model y that no longer depends on w .

In case you can't rely on the result of question (d), use the following expression for $w = \sum_{n=1}^N f(a_n, \hat{a}_n, \mu_n, \hat{\mu}_n) \phi(x_n)$, where $f(a_n, \hat{a}_n, \mu_n, \hat{\mu}_n)$ is introduced to indicate that w may still depend on the Lagrange multipliers.

- 2p **14h** Consider the following. You have access to a solver that gives you the optimizer of the dual problem, and thereby the optimal solution for the primal problem as well. It turns out that you have a hard time choosing an appropriate set of basis functions to create the feature vectors $\phi(x)$ that are used in the predictive model. What can you do to circumvent the problem of having to make specific choices for ϕ ? Explicitly state what changes you make to the predictive model and/or the optimization steps.



This page is left blank intentionally

